

AI Model Performance in CV Generation and Consistency Testing: A Practical Six-Dimension Evaluation Framework

Sai Dumpala
Birmingham City University, UK
dumpalasairamkrishnareddy@gmail.com

May 2026

Abstract

This paper describes a practical framework for evaluating large language model (LLM) output quality across six dimensions: Completeness, Contextual Utilisation, Factual Accuracy, Confidence Calibration, Consistency, and Specification Adherence. The framework came out of a real problem I noticed – LLMs producing CVs that looked finished but were missing information, or presenting wrong answers with complete confidence. I tested four models (DeepSeek, ChatGPT, Claude, and Blackbox AI) across three experiments: a basic CV generation task, a prompt wording comparison test, and a summary consistency test. The results showed that Claude performed best overall, averaging 4.7 out of 5 across all experiments. DeepSeek hallucinated significantly when instructed not to leave sections blank, inventing fake projects and contact details. Blackbox AI could not handle multiple files simultaneously. ChatGPT was honest about its limitations but inconsistent in summarisation tasks. The most important finding across all experiments was that confidence calibration failed consistently – no model ever warned users that output might be incomplete, incorrect, or different in another session. All experiment data, prompts, and notes are available at <https://github.com/SAIDUMPALA01/llm-evaluation-framework>.

Contents

1	Introduction	3
2	The Six-Dimension Framework	3
3	Experiment 1: CV Generation from a Casual Prompt	5
3.1	Setup	5
3.2	Results	5
3.3	What Actually Happened	5
3.4	Takeaway from Experiment 1	6
4	Experiment 2: Prompt Wording Comparison	6
4.1	Setup	6
4.2	Prompt 1 Results – The Safe Version	7
4.3	Prompt 2 Results – The Dangerous Version	7
4.4	Prompt 3 Results – The Vague Version	7
4.5	Key Lesson from Experiment 2	8
5	Experiment 3: Summary Consistency Test	8
5.1	Setup	8
5.2	The Three Outputs	8
5.3	Comparison	9
5.4	Scores for Experiment 3	10
6	Discussion	10
6.1	Overall Rankings	10
6.2	Completeness and Hallucination Are Connected	10
6.3	The MISSING Instruction Works	11
6.4	Consistency Is Task-Dependent	11
6.5	Confidence Calibration Is the Persistent Problem	11
7	Limitations	11
8	Conclusion	12

1 Introduction

I started thinking about this research because of something that happened to a friend of mine. He was applying for graduate jobs and used an LLM to generate his CV from his university transcripts. The model returned a document that looked professional – proper headings, clean formatting, all the right sections. He was happy with it and sent it out to several employers.

A few days later I looked at it properly. The Education section was completely empty. Not wrong – empty. The model had been given his full academic transcript including his university name, degree classification, and module grades. It just did not use any of that information. It had formatted an empty shell and called it a finished CV.

What bothered me even more was what the model said when it returned the output. It said something like “Here is your professional CV, ready to send.” No warning. No mention of missing content. My friend had no way of knowing anything was wrong just by looking at it.

That experience made me think about how we actually measure whether an LLM has done a good job. The standard approaches did not seem good enough to me. BLEU scores measure how similar output is to a reference translation [1]. Pass/fail tests check whether code executes [3]. ROUGE measures summarisation overlap [2]. These are all useful for narrow tasks, but none of them answer the questions that matter most in everyday use: Did the model use the information I gave it? Did it follow my formatting instructions? Does it know when it is missing something?

Hallucination has been studied quite a lot [4] and newer models have improved on it [5], but it is still present. What has received less attention is what I would call the invisible failure – outputs that are incomplete or misleading but presented with the same confidence as correct ones. The user has no signal to look more carefully.

More recent evaluation work has moved toward multi-dimensional frameworks. Yu et al. [6] developed DeCE, which separates precision and recall in LLM evaluation. Lauer et al. [7] evaluated LLMs on business process modelling. Sharma et al. [8] focused on health queries in Nepali. Wu et al. [9] built automated literature survey tools. These show growing awareness that a single metric is not enough. But none of them specifically address the user experience of providing personal information and expecting it to be extracted correctly.

That is the gap I tried to fill. I built a six-dimension framework from real observations of LLM failures, then tested it across three practical experiments. This paper describes the framework, the experiments, and what I found.

2 The Six-Dimension Framework

I did not start with a theory and then look for evidence. I started with broken outputs and worked backwards. Over several weeks of using LLMs for real tasks – generating documents, summarising content, answering questions – I kept notes every time something went wrong. Four patterns appeared again and again.

The first pattern was outputs that looked complete but were not. The model produced nicely

structured documents with all the right headings, but several sections had placeholder text or nothing at all. Because the formatting was correct, the problem was easy to miss.

The second pattern was consistent overconfidence. Not once in my early testing did a model say it was unsure about something, or that it could not find a piece of information I had given it. Every output was presented as finished and correct, regardless of what was actually in it.

The third pattern was ignoring provided context. Even when I gave a model a full document and explicitly asked it to use that information, it sometimes wrote generic content instead. The information was there. The model just did not extract it.

The fourth pattern was inconsistency. The same prompt in different sessions produced meaningfully different outputs – not just different wording, but different content, different completeness, different emphasis.

From these four patterns I defined the six dimensions shown in Table 1.

Table 1: Six quality dimensions for LLM evaluation

Dimension	What it means
Completeness	Does the output include everything the user asked for? Missing sections or placeholder text count as failures.
Contextual Utilisation	Did the model actually use the information provided? This is different from hallucination. Hallucination is inventing things. This is ignoring what you gave.
Factual Accuracy	Is the output correct? No invented degrees, fake projects, or wrong dates.
Confidence Calibration	Does the model acknowledge when it is uncertain or when information is missing? Or does it always present output as complete and correct?
Consistency	Does the same prompt in separate sessions produce equivalent quality output? Variation in facts or completeness counts as a failure.
Specification Adherence	Did the model follow the formatting and structural rules given in the prompt?

I scored each dimension from 1 to 5. A score of 5 means no problems at all. A score of 4 means minor issues only. A score of 3 means noticeable problems but the output is still partly usable. A score of 2 means the dimension mostly failed. A score of 1 means complete failure. All scoring was done by hand, reviewing each output carefully against the input I had provided.

One thing I want to be clear about: scoring by hand introduces subjectivity. I was the only scorer. A future version of this framework should involve multiple independent scorers to reduce that bias.

3 Experiment 1: CV Generation from a Casual Prompt

3.1 Setup

I tested four models: DeepSeek, ChatGPT, Blackbox AI (MiniMax M2.7), and Claude Sonnet 4.6. Each received the same input – my actual academic and personal information – and the same prompt. I ran all tests on 6 May 2026 using Chrome on the same device.

The prompt I used was exactly as I typed it, with no editing:

i want resume for computer science job i mean softerware engineering role like entry level i attech my bcs degree and master degree and my master final year i did a project medical image processing and my linkedin link is there www.linkedin.com/in/sai-dumpala and github link is there <https://github.com/SAIDUMPALA01> too my number +447761001792 addrees 13, norkflok rd gravesend da122rx, uk and my gmail id: dumpalasairamkrishnareddy@gmail.com give me pdf type

I attached both of my real degree transcripts. I deliberately used a casual, unedited prompt to reflect how real users actually type requests – not how a researcher would phrase a test prompt.

Before running the tests I thought all four models would give something usable. I was curious which one would use my Medical Image Processing project properly and which one would hallucinate least.

I was partially right and partially wrong.

3.2 Results

Response times were similar across all models: DeepSeek took 6–10 seconds, ChatGPT 5–9 seconds, Blackbox 6–10 seconds, and Claude 5–8 seconds. None of the response times were a practical issue.

Table 2 shows my scores for each model.

Table 2: Experiment 1 scores (out of 5)

Model	Completeness	Contextual Util.	Factual Acc.	Conf. Calibration
DeepSeek	3.5	4.0	2.5	2.0
Claude AI	4.8	5.0	5.0	4.0
Blackbox AI	2.5	3.0	3.0	2.5
ChatGPT	4.0	4.0	4.0	3.5

3.3 What Actually Happened

DeepSeek returned a nicely formatted, ATS-friendly resume within about 8 seconds. It used my LinkedIn and GitHub links correctly. It described my Medical Image Processing project in reasonable detail, mentioning segmentation and classification algorithms. But then it added two projects I never mentioned: an “E-Commerce Backend Simulation” and a “Personal Portfolio & Blog”. I did not ask for fake projects. These were inventions. So factual accuracy scored 2.5.

It also used placeholder university names like “[University Name, UK]” even though my actual university – Birmingham City University – was clearly visible in the attached transcript.

Claude Sonnet 4.6 was the best. It somehow extracted my actual university names from the transcripts – Birmingham City University for my MSc and Rayalaseema University for my BSc. It listed real module names, included internships from my BSc years, and noted I am eligible to work in the UK. The PDF output looked professional. I gave it 4.8 – would have been 5 if it had included my phone number without me needing to point it out. That is a minor complaint. Overall, it was the only model that felt like it had actually read what I gave it.

Blackbox AI was disappointing in a specific way. It could only process one PDF at a time, and I had two transcripts (BSc and MSc). It read only the BSc transcript and ignored the MSc entirely. My Master’s degree was completely missing from the output. It also said the file was unreadable in some attempts. For any task involving multiple documents, Blackbox is not ready.

ChatGPT gave a different kind of failure. It warned me that my uploaded transcripts had “expired” and could not be read. This was strange because the PDFs were fresh. As a result it produced a generic CV with no real grades, no module names, and no university details. But here is the thing – it was honest about the failure. It told me clearly what had gone wrong and offered to try again if I re-uploaded the files. That honesty is why its confidence calibration score (3.5) is higher than DeepSeek’s (2.0), even though the output was less complete. DeepSeek produced a more detailed CV but said it was professional without flagging the invented projects. ChatGPT produced a worse CV but told me why.

3.4 Takeaway from Experiment 1

Claude won clearly. DeepSeek produces good-looking output but invents content when information is missing. Blackbox cannot handle multi-file tasks. ChatGPT is cautious, which costs detail but gains honesty.

4 Experiment 2: Prompt Wording Comparison

4.1 Setup

After Experiment 1 I wanted to understand how much the phrasing of a prompt changes the output quality, particularly around hallucination. I used the same four models and the same transcripts but tested three different prompts.

- **Prompt 1 (strict + MISSING instruction):** *“You are an expert CV writer. Use ONLY the information I provide below. Fill every section. If you cannot find information for a section, write MISSING instead of leaving it blank or guessing.”*
- **Prompt 2 (fill everything):** *“Create a professional CV from this information. Fill in ALL sections completely. Do not leave anything blank.”*
- **Prompt 3 (vague):** *“Make me a CV from this information.”*

I expected Prompt 1 to produce the most honest output and Prompt 2 to be the most dangerous.

4.2 Prompt 1 Results – The Safe Version

Every model followed the MISSING instruction. No model invented content. DeepSeek wrote MISSING for contact fields it could not find, which was exactly the right behaviour. Claude filled all known information and wrote MISSING for unknowns – it also correctly extracted module names from the transcripts. ChatGPT asked me to provide information directly rather than extracting from PDFs, which was cautious but unhelpful. Blackbox reported it could not read the file.

Table 3: Experiment 2, Prompt 1 scores

Model	Completeness	Factual Acc.	Hallucination	Overall
DeepSeek	4.0	5.0	None	3.5
ChatGPT	3.5	5.0	None	3.8
Blackbox	2.5	5.0	None	2.8
Claude	4.5	5.0	None	4.5

The lesson from Prompt 1 is that the MISSING instruction works. All models followed it. No one invented data.

4.3 Prompt 2 Results – The Dangerous Version

This is where things got interesting. When told to fill ALL sections completely without any instruction about what to do with missing information, DeepSeek hallucinated significantly. It invented a phone number with a +91 Indian prefix (I am in the UK). It created a fake email address and a fake GitHub profile URL. It added three completely fabricated projects – an “Intelligent Task Scheduler,” a “Student Performance Analytics” project, and a “Digital Literacy Drive.” None of these exist. It also listed a university professor as a reference with no basis for doing so.

DeepSeek interpreted “do not leave anything blank” as permission to invent. That is the wrong response. The correct response, which Claude gave, was to use placeholders like “[Your phone number]” with a clear note explaining that the information was not available in the provided documents.

Table 4: Experiment 2, Prompt 2 scores

Model	Hallucination	Fabricated Items	Overall
DeepSeek	Yes	Phone, email, GitHub, 3 projects	3.4
ChatGPT	No	None	4.2
Blackbox	No	None (file error)	3.0
Claude	No	None (used placeholders)	4.8

4.4 Prompt 3 Results – The Vague Version

The minimal prompt produced results similar to Prompt 1 across all models. Claude scored approximately 4.6, DeepSeek 3.9, ChatGPT 4.0, and Blackbox 3.2. Interestingly, DeepSeek did

not hallucinate on the vague prompt the way it did on Prompt 2. This suggests the hallucination was specifically triggered by the “do not leave anything blank” instruction rather than being a general tendency.

4.5 Key Lesson from Experiment 2

Never tell an LLM to “not leave anything blank” unless you have given it every single piece of information. That phrase is a hallucination trigger – at least for DeepSeek. The safer instruction is “write MISSING for any information you cannot find.” That worked across all models and produced honest, usable output.

Claude seemed more resistant to prompt-induced hallucination than DeepSeek. Why this is the case is not something I can say for certain from my testing – it may relate to fine-tuning differences – but the practical implication is that prompt choice matters more with some models than others.

5 Experiment 3: Summary Consistency Test

5.1 Setup

The third experiment focused specifically on the Consistency dimension. I used ChatGPT (free version) and asked it to summarise the same paragraph three times in completely separate sessions with no conversation history between them. I ran all three sessions on 7 May 2026.

The paragraph I used was:

“Artificial intelligence is changing how people work. Companies use AI to automate tasks that humans used to do. AI can now write emails, analyse data, and create reports. This automation has increased productivity in many industries. However, there are concerns about job displacement and the reliability of AI-generated content. Experts disagree about whether the long-term impact will be mostly positive or negative.”

The prompt was identical in each session: *“Summarise this paragraph in exactly 5 bullet points.”*

Before running the experiment I thought the outputs would be almost identical – maybe some differences in word order. I was wrong.

5.2 The Three Outputs

Session 1 (started 21:06):

- AI is changing how people work.
- Companies use AI to automate human tasks.
- AI can write emails, analyse data, and create reports.
- Automation has increased productivity in many industries.
- Concerns exist about job losses, content reliability, and long-term effects.

Session 2:

- AI is changing how people work.
- Companies use AI to automate tasks (e.g., writing emails, analysing data, creating reports).
- Automation has boosted productivity across many industries.
- There are concerns about job displacement and the reliability of AI content.
- Experts have different opinions on AI’s long-term impact.

Session 3:

- AI is transforming work patterns.
- AI automates tasks like email writing, data analysis, and report generation.
- Productivity has increased due to automation.
- Concerns include job losses and unreliable AI content.
- Experts disagree about AI’s long-term effects.

5.3 Comparison

Table 5: Experiment 3 comparison across three sessions

Key Point	Session 1	Session 2	Session 3
AI changing work	Yes	Yes	Yes
Automation examples	Two bullets (split)	One bullet	One bullet
Productivity	Yes	Yes	Yes
Job concerns	Merged with Point D	Separate bullet	Separate bullet
Expert disagreement	Weakened ("long-term effects")	Clearly stated	Clearly stated
Bullet count	5	5	5

All three sessions gave exactly 5 bullets as instructed and used only content from the paragraph. On those dimensions – Specification Adherence and Contextual Utilisation – the model performed well.

The problem was structural. Session 1 split the automation examples into two separate bullets, which left only one bullet to cover both job displacement concerns and expert disagreement. As a result, expert disagreement was weakened to “long-term effects” – a vague phrase that does not capture the fact that experts actually hold different views.

Sessions 2 and 3 handled this differently. Both kept the automation examples in a single bullet, which freed up space for separate treatment of job concerns and expert disagreement. Both sessions clearly stated that experts disagree, which is a more accurate representation of the paragraph.

This structural difference has a real consequence. A person reading only Session 1 would think the paragraph is expressing general concern about AI’s future. A person reading Session 2 or 3 would

understand that experts actively and specifically disagree. These are different understandings of the same text. And in every case, ChatGPT presented its summary as the definitive correct answer.

5.4 Scores for Experiment 3

Table 6: Experiment 3 scores per session

Dimension	Session 1	Session 2	Session 3	Average
Completeness	4	5	5	4.67
Contextual Utilisation	5	5	5	5.00
Factual Accuracy	4	5	5	4.67
Confidence Calibration	2	2	2	2.00
Consistency (overall)		3/5		3.00
Specification Adherence	5	5	5	5.00

Confidence Calibration scored 2 across all three sessions. In none of them did ChatGPT say “this is one possible summary” or “results may vary if you ask again.” Each session was presented as final. A user would never know that Session 1 expressed expert disagreement less clearly than Sessions 2 and 3 unless they ran the test multiple times themselves. Most users do not do that.

If I were a student using ChatGPT to summarise a paper for a literature review and I only ran it once – getting Session 1 – I would miss the fact that experts disagree. That is a real mistake that would affect my understanding of the source material.

6 Discussion

6.1 Overall Rankings

Averaging scores across all experiments gives the following ranking:

1. **Claude AI – 4.7/5**: Consistent, no hallucinations across any prompt, good information extraction from transcripts.
2. **ChatGPT – 4.0/5**: Honest about failures but inconsistent in summarisation. Better on document tasks than consistency tasks.
3. **DeepSeek – 3.6/5**: Strong structure and formatting, but hallucinates when prompted to fill everything completely.
4. **Blackbox AI – 3.0/5**: File handling limitations make it unreliable for any task involving multiple documents.

6.2 Completeness and Hallucination Are Connected

The most important finding from Experiments 1 and 2 is that completeness failures and hallucination are connected in a specific way. When a model is told to produce a complete document without any instruction about how to handle missing information, some models respond by inventing content rather than admitting they do not have it.

DeepSeek’s behaviour on Prompt 2 is the clearest example. It treated “fill everything” as permission to fabricate. A user reading that CV would have no way of knowing the phone number was invented, the projects were fake, and the references were made up. The output looked completely real.

6.3 The MISSING Instruction Works

One practical and useful finding from this study is that telling a model to write MISSING for unknown information significantly improves confidence calibration. All four models followed this instruction when given it. None of them invented content when given the MISSING instruction.

This is good news because it means the overconfidence problem is partly addressable through better prompting. But it should not have to be. Users should not need to know about this technique to get honest output. It should be the default behaviour.

6.4 Consistency Is Task-Dependent

Experiment 3 showed that consistency issues are not limited to document generation tasks. ChatGPT produced meaningfully different summaries of the same paragraph across three sessions. Sessions 2 and 3 were consistent with each other, but Session 1 was structurally different in a way that affected the accuracy of the takeaway.

This finding extends the consistency concern beyond what has been shown in previous studies on factual Q&A. Even in a constrained summarisation task with a fixed input, structural choices can vary between sessions in ways that change what the user understands.

6.5 Confidence Calibration Is the Persistent Problem

Across all three experiments and all four models, confidence calibration was the most consistently poor dimension. No model in any experiment expressed uncertainty, flagged potential inaccuracies, or warned that results might vary. DeepSeek said “Here is your professional CV” after inventing three fake projects. ChatGPT presented each of its three summaries as the definitive correct answer. Claude said “Your CV is ready” even when contact details were missing.

This is the dimension that matters most for real users because it is the one that prevents them from knowing when to check carefully. If a model flagged uncertainty, users would know to verify. Because no model does, users trust output that may be incomplete, inaccurate, or different from what they would have got if they asked again.

7 Limitations

There are several limitations I want to acknowledge honestly.

First, all scoring was done by me alone. There was no second reviewer to check my judgements. This introduces the possibility of personal bias, particularly for dimensions like Contextual Utilisation where the difference between good and acceptable performance involves judgement calls.

Second, I only tested one document generation task (CV from transcripts). The findings about DeepSeek’s hallucination behaviour may not generalise to other document types. And the results from Experiment 3 are based on a single paragraph – different paragraph types might produce different consistency patterns.

Third, I used free versions of all models. Paid versions may have higher context limits, better file handling, or different behaviours. Blackbox AI in particular might perform better with a paid plan that supports multiple file uploads.

Fourth, LLMs are updated frequently. The specific behaviours I observed in May 2026 may change in future model versions. This paper describes what I found at a specific point in time, not a permanent characterisation of any model.

Fifth, I did not test non-English tasks. All prompts and documents were in English. Results might differ for other languages.

8 Conclusion

This paper started from a practical problem: an LLM produced a CV that looked finished but was missing real information, and gave no warning about it. From that problem I built a six-dimension framework and tested it across three experiments involving four models.

The main findings are:

Claude performed best overall, with strong contextual utilisation and no hallucinations across all three prompt styles. It was the only model that consistently extracted real information from provided documents.

DeepSeek produced well-structured output but hallucinated significantly when instructed not to leave sections blank. It invented phone numbers, email addresses, GitHub profiles, and entire projects. Users who do not know to check for this are at risk.

Blackbox AI could not handle multiple file uploads, which severely limited its usefulness for real document tasks.

ChatGPT was honest about its limitations in Experiment 1 but showed consistency failures in Experiment 3, producing a structurally different summary in Session 1 that weakened a key point compared to Sessions 2 and 3.

The most important practical lesson from this work is about prompt phrasing. The instruction “write MISSING for any information you cannot find” produced honest, accurate output from all models. The instruction “do not leave anything blank” produced hallucinations from DeepSeek. That difference in one phrase changed the output from honest to fabricated. Users should know this.

The most important finding for researchers and developers is about confidence calibration. Not one of the four models, across any of the three experiments, ever expressed uncertainty, acknowledged that output might be incomplete, or warned that answers might vary across sessions. This is

the dimension that needs the most work. Until models learn to say “I am not sure” or “I could not find this information” by default, users are on their own when it comes to checking output quality.

Future work should extend this framework to other task types, involve multiple independent scorers, test paid versions of the models, and develop automated tools for measuring each dimension at scale.

All data, prompts, and raw session transcripts are available at: <https://github.com/SAIDUMPALA01/llm-evaluation-framework>

Data and Code Availability

All experiment data, raw prompts, session transcripts, and scoring sheets are publicly available on GitHub: <https://github.com/SAIDUMPALA01/llm-evaluation-framework>

References

- [1] K. Papineni, S. Roukos, T. Ward, and W. Zhu, “BLEU: a method for automatic evaluation of machine translation,” in *Proc. 40th Annual Meeting of the Association for Computational Linguistics (ACL)*, pp. 311–318, 2002.
- [2] C. Y. Lin, “ROUGE: A package for automatic evaluation of summaries,” in *Text Summarization Branches Out: Proc. ACL-04 Workshop*, pp. 74–81, 2004.
- [3] M. Chen et al., “Evaluating large language models trained on code,” *arXiv preprint arXiv:2107.03374*, 2021.
- [4] Z. Ji et al., “Survey of hallucination in natural language generation,” *ACM Computing Surveys*, vol. 55, no. 12, pp. 1–38, 2023.
- [5] OpenAI, “GPT-4 technical report,” *arXiv preprint arXiv:2303.08774*, 2023.
- [6] F. Yu, N. Seedat, D. Herrmannova, F. Schilder, and J. R. Schwarz, “Beyond pointwise scores: Decomposed criteria-based evaluation of LLM responses,” in *Proc. EMNLP 2025 Industry Track*, 2025.
- [7] C. Lauer, P. Pfeiffer, A. Rombach, and N. Mehdiyev, “Assessing the business process modeling competences of large language models,” *arXiv preprint*, 2026.
- [8] M. Sharma et al., “Evaluating large language models responses to sexual and reproductive health queries in Nepali,” *arXiv preprint*, 2026.
- [9] S. Wu et al., “AutoSurvey2: Empowering researchers with next-level automated literature surveys,” *arXiv preprint arXiv:2510.26012*, 2025.